

CASA 4.3 MPI Parallel Processing  
Framework

Date: 2014-05-20  
Status: Development  
Page: 1 of 34

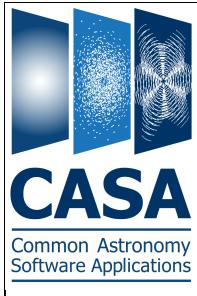
# **CASA 4.3 Parallel Processing Framework**

## **- Installation and advance user guide -**

Version: 1.0

Status: Development

<b>Prepared By:</b>		
<b>Name</b>	<b>Organisation</b>	<b>Date</b>
Justo Gonzalez	ESO	2014-05-20

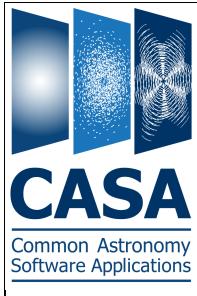


## CASA 4.3 MPI Parallel Processing Framework

Date: 2014-05-20  
Status: Development  
Page: 2 of 34

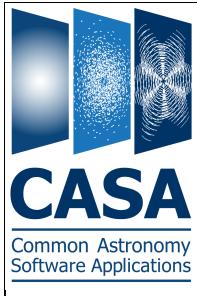
### Change Record

Version	Date	Affected Section	Change Request #	Remarks
1.0	2014-05-20	All		Initial version



## Table of Contents

<b>1 LIST OF TERMS AND ACRONYMS.....</b>	<b>4</b>
<b>2 LIST OF RELATED JIRA ISSUES.....</b>	<b>4</b>
<b>3 INTRODUCTION.....</b>	<b>5</b>
<b>4 INSTALLATION GUIDE.....</b>	<b>6</b>
4.1 OPEN MPI.....	6
4.2 MPI4PY .....	8
4.3 BOOST MPI .....	10
4.4 COMPILE CASA USING MPI COMPILERS.....	10
<b>5 MPI4CASA.....</b>	<b>11</b>
5.1 MODULE STRUCTURE.....	11
5.2 MPI4CASA INITIALIZATION .....	12
5.3 MPICOMMANDCLIENT LIFE CYCLE.....	14
5.4 MPICOMMANDCLIENT USAGE .....	16
5.5 MPICOMMANDCLIENT/SERVER LOGGING .....	21
5.6 MPICOMMANDCLIENT/SERVER ERROR HANDLING.....	22
5.7 MPICOMMANDCLIENT BASIC PYTHON EXAMPLES.....	23
5.1 MPICOMMANDCLIENT CASA TASKS EXAMPLES .....	30
5.2 MPI4CASA UNIT TEST SUIT.....	34

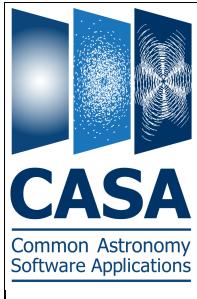


## 1 List of terms and acronyms

CASA	Common Astronomy Software Applications
MS	Measurement Set
MSs	Group of Measurement Sets
MMS	Multi Measurement Set
MMSs	Group of Multi Measurement Sets
Sub-MS	Sub-Measurement Set (component of a Multi Measurement Set)
Sub-MSSs	Group of Sub-Measurement Sets
API	Application Programming Interface
MPI	Message Passing Interface
GUI	Graphical User Interface
HPC	High Performance Computing
stl	C++ Standard Template Library
iPython	Python command shell with advanced interactivity features such as rich history, enhanced introspection, etc
InfiniBand	Network communication link used in HPC applications

## 2 List of related Jira issues

<a href="#">CAS-5795</a>	MPI-based CASA cluster infrastructure (container ticket, see issues under it)
<a href="#">CAS-5797</a>	Test integration of MPI-based ipcluster with Torque
<a href="#">CAS-5798</a>	Compare the various packages offering python bindings for MPI
<a href="#">CAS-5800</a>	Analysis of schemes to support MPI in the CASA cluster environment
<a href="#">CAS-5801</a>	Analysis of MPI implementations (system wise, apart from the python bindings)



### 3 Introduction

This is a technical document (advance user and developer guide) describing how to install and set up the necessary packages to use the new MPI-based parallel processing framework for CASA 4.3. It also contains several examples on how to use `mpi4casa`, the new Python module that enables MPI-based parallelism at the task level. The following list described the package used, and why it is needed:

1. **Open MPI:** One of the advantages of MPI is the possibility to switch from one implementation to another w/o suffering API changes. Despite of this I believe that CASA should provide a preferred MPI implementation together with the release package, and we should also have a standardized development environment, and that means selecting a group of MPI implementations and use them for development and testing. Having said this I've chosen Open MPI for my own development and testing because it is the one of the newest implementations, with extensive development and aims to support all common interconnects including InfiniBand and full MPI-2 compliance. Therefore I suggest to go ahead with it, and study in parallel what advantages / disadvantages come along with other implementations.
2. **mpi4py:** This is basically the MPI Python bindings that allow to use MPI from the Python layer. There are several choices, but `mpi4py` is the best in terms of performance, completeness and supports communication of arbitrary python objects, thus w/o requiring any serialization. It is necessary because in general CASA is best parallelized from the task level, which is the highest level possible, and thus minimizes the parallelization overhead. This is so for all the tasks that support trivial parallelism, and therefore don't need efficient communication at the C++ level. The exception is imaging, which due to the major/minor cycle structure requires efficient communication among the parallel processes.
3. **boostmpi:** This is the preferred package providing a C++ interface for MPI. It supports direct communication of all the stl library types and containers w/o requiring further serialization, thus is a really good candidate. Furthermore boost has become a standard in itself in the C++ modern development world, and it is a force pushing for many changes in the new C++ standards (e.g. many new features of C++11 are already available in C++99 via boost).
4. **mpi4casa:** This is the new CASA module implementing a complete MPI-based parallel processing framework. It is built on top of `mpi4py`, and thus honors its name. It uses a Client/Server model, where two completely different CASA environments are loaded depending on the role. The 'main' Client process is basically a normal CASA session, including all the associated GUIs and processes for interactive usage. The other Server processes load the minimal CASA environment w/o GUIs, iPython or any interactivity component, and communicate with the 'main' Client process via MPI.



## 4 Installation guide

### 4.1 Open MPI

#### 4.1.1. Download Open MPI:

```
wget http://www.open-mpi.org/software/ompi/v1.6/downloads/openmpi-1.6.5.tar.bz2
```

#### 4.1.2. Unzip and untar package

```
bunzip2 openmpi-1.6.5.tar.bz2  
tar -xvf openmpi-1.6.5.tar  
cd openmpi-1.6.5
```

#### 4.1.3. Configure enabling thread-safe support, compile and install

**NOTE:** It is absolutely necessary that you explicitly add the multi-threading support options, because they are not added by default.

```
./configure --enable-mpi-thread-multiple --enable-opal-multi-threads  
make all  
make install
```

#### 4.1.4. Update the \$HOME/.bash\_profile file to include the location of the MPI libraries in LD\_LIBRARY\_PATH

```
vim $HOME/.bash_profile  
export LD_LIBRARY_PATH=/usr/local/lib
```

#### 4.1.5. Make sure that MPI compilers and launchers are available in the PATH:

```
MPI C compiler:      /usr/local/bin/mpicc  
MPI C++ compiler:    /usr/local/bin/mpicxx  
MPI F77 compiler:    /usr/local/bin/mpif77  
MPI F90 compiler:    /usr/local/bin/mpif90  
Pre MPI 2.0 launcher: /usr/local/bin/mpirun  
MPI 2.0 launcher     /usr/local/bin/mpieexec
```

#### 4.1.6. Enable password free ssh access to all the machines running MPI:

```
ssh-keygen -t rsa  
$HOME/.ssh/id_rsa.pub | ssh user@hostname 'cat >> $HOME/.ssh/authorized_keys'
```

#### 4.1.7. Make a basic test in the localhost to make sure MPI is properly installed

```
mpirun -n 4 ls
```



4.1.8. Prepare a simple hostfile with 2 nodes and 2 slots per node following the Open MPI hostfile syntax:

```
# This is an example hostfile.  Comments begin with #
# The following node is a single processor machine:
foo.example.com
# The following node is a dual-processor machine:
bar.example.com slots=2
# The following node is a quad-processor machine, and
# we absolutely want to disallow over-subscribing it:
yow.example.com slots=4 max-slots=4
```

4.1.9. Make a test using the hostfile from the previous step:

```
[testhpc2@almahpc02 MPI]$ cat hostfile
almahpc02.ads.eso.org slots=2
almahpc01.hq.eso.org slots=2
[testhpc2@almahpc02 MPI]$ mpirun -n 4 -hostfile hostfile hostname
almahpc02
almahpc02
almahpc01
almahpc01
```



## 4.2 mpi4py

### 4.2.1. Download mpi4py

```
wget http://mpi4py.googlecode.com/files/mpi4py-1.3.1.tar.gz
```

### 4.2.2. Unzip and untar package

```
tar -xzvf mpi4py-1.3.1.tar.gz
```

### 4.2.3. Build and install the mpi4py package using the python version shipped with CASA. It will automatically detect the MPI implementation and use it for compilation.

```
/usr/lib64/casa/01/bin/python setup.py build  
/usr/lib64/casa/01/bin/python setup.py install
```

### 4.2.4. Build the also the Python interpreter with MPI enabled on it (it will create a python2.7-mpi executable)

```
/usr/lib64/casa/01/bin/python setup.py build_exe  
/usr/lib64/casa/01/bin/python setup.py install_exe
```

### 4.2.5. Create a link so that 'python' points to python2.7-mpi:

```
cd /usr/lib64/casa/01/bin  
ln -s python2.7-mpi python
```

### 4.2.6. Run a simple parallel python test using one xterm per python process:

```
mpirun -n 2 -xterm 0,1 /usr/lib64/casa/01/bin/python
```



4.2.7. Copy and paste the following test code in both xterm python terminals. It basically sends a dictionary from the rank 0 python process to the rank 1 python process. Make sure that the dictionary is received in the rank 1 process by printing it.

```
from mpi4py import MPI

comm = MPI.COMM_WORLD
rank = comm.Get_rank()

if rank == 0:
    data = {'a': 7, 'b': 3.14}
    comm.send(data, dest=1, tag=11)
elif rank == 1:
    data = comm.recv(source=0, tag=11)
```

The image shows two terminal windows side-by-side. The left window is titled "Rank 0 (on almahpc02)" and the right window is titled "Rank 1 (on almahpc02)". Both windows show Python 2.7.5 running on a Linux system. The code in both windows is identical, demonstrating MPI communication between ranks 0 and 1. In the Rank 0 window, the output shows the data being sent from rank 0 to rank 1. In the Rank 1 window, the output shows the data being received by rank 1.

```
Rank 0 (on almahpc02)
Python 2.7.5 (default, Sep 26 2013, 14:33:04)
[GCC 4.1.2 20080704 (Red Hat 4.1.2-54)] on linux2
Type "help", "copyright", "credits" or "license" for more information.
>>> from mpi4py import MPI
>>>
>>> comm = MPI.COMM_WORLD
>>> rank = comm.Get_rank()
>>>
>>> if rank == 0:
...     data = {'a': 7, 'b': 3.14}
...     comm.send(data, dest=1, tag=11)
... elif rank == 1:
...     data = comm.recv(source=0, tag=11)
...
>>> print rank
0
>>> print data
{'a': 7, 'b': 3.14}
>>> 
```

```
Rank 1 (on almahpc02)
Python 2.7.5 (default, Sep 26 2013, 14:33:04)
[GCC 4.1.2 20080704 (Red Hat 4.1.2-54)] on linux2
Type "help", "copyright", "credits" or "license" for more information.
>>> from mpi4py import MPI
>>>
>>> comm = MPI.COMM_WORLD
>>> rank = comm.Get_rank()
>>>
>>> if rank == 0:
...     data = {'a': 7, 'b': 3.14}
...     comm.send(data, dest=1, tag=11)
... elif rank == 1:
...     data = comm.recv(source=0, tag=11)
...
>>> print rank
1
>>> 
```



## 4.3 Boost MPI

### 4.3.1. Install bjam (boost setup utility)

```
wget http://sourceforge.net/projects/boost/files/boost-jam/3.1.18/boost-jam-3.1.18.tgz/download
tar -xzvf boost-jam-3.1.18.tgz
cd boost-jam-3.1.18
./build.sh
cp bin.linuxx86_64/bjam /usr/local/bin/
```

### 4.3.2. Download and unpack boost:

```
wget
http://sourceforge.net/projects/boost/files/boost/1.41.0/boost_1_41_0.tar.gz/download
tar -xzvf boost_1_41_0.tar.gz
cd boost_1_41_0
```

### 4.3.3. Edit the boost configuration file to specify the python binary and add the boost::mpi module

```
vim ./tools/build/v2/user-config.jam
=> Add using python : 2.7 : /usr/lib64/casa/01/bin/python ;
=> Add using mpi ;
```

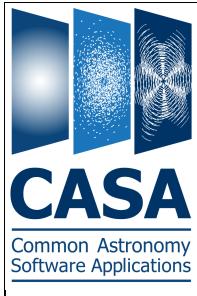
### 4.3.4. Using bjam configure boost specifying as prefix the location where the CASA boost packages is installed (/usr/lib64/casa/01)

```
bjam --debug-configuration --prefix= /usr/lib64/casa/01 install
```

## 4.4 Compiling CASA using MPI compilers

### 4.4.1. Now all the pieces are ready to compile CASA C++ code with full MPI support. To do it so it only necessary to specify the MPI compilers in the cmake configuration step.

```
cmake -DCMAKE_Fortran_COMPILER=/usr/bin/mpif90 -DCMAKE_C_COMPILER=/usr/bin/mpicc -
DCMAKE_CXX_COMPILER=/usr/bin/mpic++ (... other options ...)
```



## 5 mpi4casa

### 5.1 Module structure

The new module to support MPI at the CASA task level is called `mpi4casa`. It is located under `gwrap/python/scripts/mpi4casa`. The following table describes the contents of the module:

<code>MPIEnvironment.py</code>	Static class in charge of Initializing MPI and properly set the MPI-related variables.
<code>MPICommunicator.py</code>	Singleton class, which centralizes all the MPI, calls, using the appropriate targets and tags.
<code>MPICommandClient.py</code>	Singleton class, which provides access to high-level methods to: - Send CASA commands to the remote side (with blocking, non-blocking modes, targeting a single or multiple servers) - Get the response from the remote calls (either return variables or stack traces in case of error)
<code>MPIMonitorClient.py</code>	Singleton class used primarily by <code>MPIMonitorClient</code> , to monitor the state of the remote servers in terms of activity (busy/idle) and responsiveness.
<code>MPICommandServer.py</code>	Counterpart for <code>MPICommandClient</code> : Singleton class with a service thread that checks for incoming MPI messages containing CASA (or in general Python) command requests. It executes them and sends back the response (return code), and back trace in case if error.
<code>MPIMonitorServer.py</code>	Counterpart for <code>MPIMonitorClient</code> : Singleton class with a service thread that checks for incoming MPI messages containing check status requests.
<code>mpi4casapy.py</code>	Runnable python script used to launch the MPI Server environment.
<code>task_wrappers.py</code>	Python script containing the import definitions for the CASA tasks wrappers to be used in the MPI Server environment.
<code>task_macros.py</code>	Python script containing the macro definitions for the CASA tasks wrappers to be used in the MPI Server environment.
<code>test_mpi4casa.py</code>	Unit test suit with more than 60 cases covering: - MPI client/server command cases: <ul style="list-style-type: none"><li>* evaluation with return code or execution</li><li>* blocking and non-blocking modes</li><li>* single and multiple targets</li><li>* parameters provided via string or dictionary</li><li>* error/exception handling</li><li>* server timeout state handling</li><li>* server timeout recovery</li></ul>



## 5.2 mpi4casa initialization

As you can see in the previous section mpi4casa uses a Client/Server execution model:

- **MPIClient:** 'Main' MPI instance with all interactivity features (GUIs, iPython shell, async, etc). That is, a complete normal CASA instance.
- **MPIServer:** Minimal CASA environment without any interactivity features
  - Basically a python script running in the background
  - CASA env. Variables are sent as a dictionary via MPI from MPIClient (e.g. log file, data paths, etc)
- **Convention:** The established convention is that the Client is the process with highest rank, for example with 4 processors, the Client process should be the process with rank 3 (in MPI rank process rank numbering starts with 0). For this reason:
  - stdin must be redirected only for the process with higher rank

```
mpirun -n 4 -stdin 3 casapy
```

- xterm must be initiated only for the process with higher rank

```
mpirun -n 4 -xterm 3 casapy
```

**NOTE:** stdin/out redirection from the terminal is not the recommended way to interact with the client due to a miss-formatting that occurs when stdin/out is redirected through the mpirun process. Instead we recommend to use a dedicated xterm window. However stdin/out redirection from terminal is the only way to interact with the client in MACOX



- **Initialization success:** After launching CASA with mpirun, if the initialization is successful you will see the following message:

```
casa::casapy::casa@almahpc02:MPIClient MPI Enabled at host almahpc02 with
rank 3 as MPIClient using MPI version 2.1 from Open MPI v1.6.5 implementation
```

The screenshot shows a terminal window titled "Rank 3 (on almahpc02)". The log output includes:

```
on: Mon, 28 Apr 2014
*** Loading ATNF ASAP Package...
*** ... ASAP (4.2.0a rev#2919) import complete ***
#####
Major interface changes to SINGLE DISH tasks have been
taken place in CASA 4.2.2 release

The interface of the following tasks are modified:
sdbaseline, sdcal, sdcal2, sdfit, sdflag, sdgrid,
sdimaging, sdmath, sdplot, sdreduce, sdsave, and sdstat.
Additionally, a new task called sdaverage is available. Task
sdsmooth has been incorporated in the new task and removed.

The tasks with old interfaces are available with name
{taskname}old and kept by CASA 4.3 release. Users are
advised to update existing scripts.

#####
2014-05-16 17:05:24     INFO    casa::casapy::casa@almahpc02:MPIClient MPI Enab
led at host almahpc02 with rank 3 as MPIClient using MPI version 2.1 from Open M
PI v1.6.5 implementation
```

This message indicates that the MPI version provided by the underlying MPI implementation is correct, and also that thread-safe support is enabled.

Also, if all the processes are deployed in the same machine, you can see via pstree how the processes are deployed from mpirun:

```
[testhpc2@almahpc02 MPI]$ ( mpirun -n 4 -xterm 3 casapy --nologger --log2term &> mpirung.log ) &
[1] 8704
[testhpc2@almahpc02 MPI]$ pstree 8704
mpirun---3*[python]
  \-xterm---python---2*[casaviewer]
    |   \-dbus-daemon
    |   \-ipcontroller
    |   \-python
    \-3*[{python}]
```

- **Initialization failure:** If the mpi4casa initialization is not sucessfull (e.g.: the underlying MPI implementation does not have thread-safe support) and error message appears and it is not possible to use instantiate MPIClient:

```
WARN    casa::casapy::casa      Provided MPI implementation (Open MPI v1.6.5)
is not thread safe configured, maximum thread safe level supported is: MPI THREAD SINGLE

WARN    casa::casapy::casa+    NOTE: In most MPI implementations thread-safety
can be enabled at pre-compile, by setting explicit thread-safe configuration options,

WARN    casa::casapy::casa+    e.g. (MPI 1.6.5) --enable-mpi-thread-multiple
```



### 5.3 MPICommandClient life cycle

As mentioned previously MPICommandClient creational pattern is a singleton. This design choice is useful to avoid multiple initializations, which can cause resource problems. Nevertheless, the user can control when the Client services are initialized or de-initialized, in order to waste CPU resources when parallelization is not necessary.

**NOTE:** mpi4py is already integrated with ParallelTaskHelper, so it is not necessary to do anything (in particular initialize MPICommandClient) in order to run tasks in parallel. ParallelTaskHelper re-uses or initializes the MPICommandClient singleton if necessary.

- **MPIClient initialization:** MPIClient is a singleton, therefore it is necessary to initialize it only once, and all the other instances created afterwards will point to the first one. This is done transparently for the user.

```
CASA <2>: from mpi4casa.MPICommandClient import MPICommandClient
CASA <3>: client = MPICommandClient()
CASA <4>: client.start_services()
```

MPIClient will not only initialize the client service, but also send a message to the servers so that they are initialized:

```
Rank 3 (on almahpc02)
#####
2014-05-16 17:58:15 INFO casa::casapy::casa@almahpc02:MPIClient MPI Enabled at host almahpc02 with rank 3 as MPIClient using MPI version 2.1 from Open MPI v1.6.5 implementation
CASA <2>: from mpi4casa,MPICommandClient import MPICommandClient
CASA <3>: client = MPICommandClient()
CASA <4>: client.start_services()
2014-05-16 17:58:21 INFO casa::MPICommandClient::send_start_service_signal::casa@almahpc02:MPIClient Sending start service signal to all servers
2014-05-16 17:58:21 INFO casa::MPICommunicator::control_service_request_send_all::casa@almahpc02:MPIClient Sending start service signal to server 0
2014-05-16 17:58:21 INFO casa::MPICommunicator::control_service_request_send_all::casa@almahpc02:MPIClient Sending start service signal to server 1
2014-05-16 17:58:21 INFO casa::MPICommunicator::control_service_request_send_all::casa@almahpc02:MPIClient Sending start service signal to server 2
2014-05-16 17:58:23 INFO casa::MPICommandClient::send_start_service_signal::casa@almahpc02:MPIClient Server with rank 0 started at almahpc02 with PID 17610
2014-05-16 17:58:23 INFO casa::MPICommandClient::send_start_service_signal::casa@almahpc02:MPIClient Server with rank 2 started at almahpc02 with PID 17612
2014-05-16 17:58:24 INFO casa::MPICommandClient::send_start_service_signal::casa@almahpc02:MPIClient Server with rank 1 started at almahpc02 with PID 17611
2014-05-16 17:58:24 INFO casa::MPICommandClient::send_start_service_signal::casa@almahpc02:MPIClient Received response from all servers to start service signal
2014-05-16 17:58:24 INFO casa::MPIMonitorClient::start_ping_status_response_handler_service::casa@almahpc02:MPIClient MPI ping status response handler service started
2014-05-16 17:58:24 INFO casa::MPIMonitorClient::start_monitor_status_service::casa@almahpc02:MPIClient MPI monitor status service started
2014-05-16 17:58:24 INFO casa::MPICommandClient::start_command_request_queue_service::casa@almahpc02:MPIClient MPI command request queue service started
2014-05-16 17:58:24 INFO casa::MPICommandClient::start_command_response_handler_service::casa@almahpc02:MPIClient MPI command response handler service started
2014-05-16 17:58:24 INFO casa::MPICommandClient::start_services::casa@almahpc02:MPIClient All services started
CASA <5>: 
```



- **MPICommandClient de-initialization:** Since MPICommandClient is a singleton its life cycle is handled carefully in a transparent way for the user, so that the singleton instance is not deleted as long as other Python objects are using it. This is not done through a `__del__` operator but using the `atexit` module instead, to register an exit function which is executed when python is finalized.

Nevertheless it is possible to stop the services manually after running all the parallel processing steps

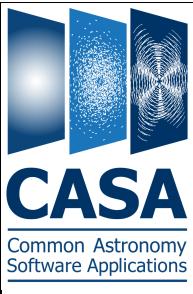
**CAUTION:** When MPIClient services are stopped, the entire parallel system is stopped, including finalization of the MPIServer processes. This means that it is not longer possible to use the parallel processing framework after stopping the MPIClient services.

```
CASA <4>: client.stop_services()
```

Otherwise, if the MPIClient services are not stopped manually, the stop function is automatically invoked when Python exits as described before:

```
Rank 3 (on almahpc02)

CASA <5>: exit()
leaving casapy...
2014-05-16 17:58:39 INFO casa::MPIMonitorClient::stop_monitor_status_service::casa@almahpc02:MPIClient    MPI monitor status service stopped
2014-05-16 17:58:40 INFO casa::MPIMonitorClient::stop_ping_status_response_handler_service::casa@almahpc02:MPIClient    MPI ping status response handler service stopped
2014-05-16 17:58:40 INFO casa::MPICommandClient::stop_command_request_queue_service::casa@almahpc02:MPIClient    MPI command request queue service stopped
2014-05-16 17:58:40 INFO casa::MPICommandClient::stop_command_response_handler_service::casa@almahpc02:MPIClient    MPI command response handler service stopped
2014-05-16 17:58:40 INFO casa::MPICommandClient::send_stop_service_signal::casa@almahpc02:MPIClient    Sending stop service signal to all servers
2014-05-16 17:58:40 INFO casa::MPICommunicator::control_service_request_send_all::casa@almahpc02:MPIClient    Sending stop service signal to server 0
2014-05-16 17:58:40 INFO casa::MPICommunicator::control_service_request_send_all::casa@almahpc02:MPIClient    Sending stop service signal to server 1
2014-05-16 17:58:40 INFO casa::MPICommunicator::control_service_request_send_all::casa@almahpc02:MPIClient    Sending stop service signal to server 2
2014-05-16 17:58:40 INFO casa::MPICommandClient::send_stop_service_signal::casa@almahpc02:MPIClient    Stop service signal sent to all servers
2014-05-16 17:58:40 INFO casa::MPICommandClient::stop_services::casa@almahpc02:MPIClient All services stopped
```



## 5.4 MPICommandClient usage

MPICommandClient has mainly two public methods to interface with the user, in order to send commands to the remote servers, and retrieve the results. Additionally there is a method to retrieve the status of the remote servers, if that is necessary for debugging reasons.

**NOTE:** mpi4py is already integrated with ParallelTaskHelper, so it is not necessary to do anything (in particular initialize MPICommandClient) in order to run tasks in parallel. ParallelTaskHelper re-uses or initializes the MPICommandClient singleton if necessary.

- `def push_command_request(self,command,block=False,target_server=None,parameters=None)`

This method allows sending CASA/Python commands to the remote servers, so that they are evaluated (with return code) or executed (w/o return code).

It is possible to call this method in block and non-block mode. When running in non-block mode it will return a list of integers, corresponding to the command request ids, which can be used in a second stage to retrieve the response via the public method 'get\_command\_response' (explained later in this section). When running in blocking mode it will return a list of dictionaries, one per command request, containing the response parameters.

It is possible to specify a defined target server or group of target servers, by providing a list of integers, corresponding to the MPI rank of the target servers.

It is also possible to specify the command parameters via an auxiliary dictionary, and the command will be executed in the remote servers after injecting the variables specified in the parameters dictionary in the python global variables namespace.



- **command:** String containing the Python/CASA command to be executed. The parameters can be specified in two ways:

- ◆ Within the command in itself, also as strings:

```
"flagdata(backup=True)"
```

- ◆ In the parameters dictionary, using the native types:

```
parameters={"backup":True}
```

- **block:** Boolean to control whether command request is executed in blocking or not:
  - ◆ Block=True: It will block until the command is queue, sent to a remote process, executed, and the corresponding response is received.
  - ◆ Block=False: It will not block, just register the command request and return an identifier which can be used later on to get the command response
- **target\_server:** List of integers (server ids) which are the target for the command. The ids correspond to the MPI process rank of the target servers

**NOTE:** Remember that according with the mpi4casa convention the process with highest rank corresponds to the client, and the other process are the servers. e.g.: In a 4 processors run, the ranks with id 0,1,2 correspond to the servers, and can be used as target\_server, and the process with rank 0 corresponds to the client.

- ◆ target\_server=None: The command will be executed by the first available server
  - ◆ target\_server=2: The command will be executed by the server n#2 as soon as it is available
  - ◆ target\_server=[0,1]: The command will be executed by the servers n #2 and #3 as soon as they are available
- **parameters** (optional): dictionary containing the parameters to be used in the command execution. These parameters will be injected in the local variable space of the server and deleted after execution. Any python (pickable) object is accepted.



- **returns:**

- ◆ In non-blocking mode: It will return immediately, giving the responses already available as a list of dictionaries, one per command request, containing the response parameters.

```
CASA <8>: client.push_command_request("a+b",False,[0],{'a':1,'b':2})
Out[8]: [3]
2014-05-19 15:10:49    INFO  casa::MPICommandClient::command_request_queue_service::casa@almahpc02:MPIClient Command request with id# 3 sent to server n# 0
```

- ◆ In blocking mode: It will not return until the responses for all the command requests have been received. Then it returns a list of dictionaries, one per command request, containing the response parameters.

```
CASA <7>: client.push_command_request("a+b",True,[0],{'a':1,'b':2})
2014-05-19 15:10:23    INFO  casa::MPICommandClient::command_request_queue_service::casa@almahpc02:MPIClient Command request with id# 2 sent to server n# 0
2014-05-19 15:10:24    INFO  casa::MPICommandClient::command_response_handler_service::casa@almahpc02:MPIClient      Command request with id 2 successfully handled by server n# 0
Out[7]:
[{'command': 'a+b',
 'id': 2,
 'mode': 'eval',
 'parameters': {'a': 1, 'b': 2},
 'ret': 3,
 'server': 0,
 'status': 'response received',
 'successful': True,
 'traceback': None}]
```

The command response dictionary contains the following keys:

keyword	type	Description
command	string	CASA/Python command that was executed
id	integer	Identifier assigned to the command request
mode	string	There are two command modes: - eval: The command was executed in eval mode which returns a variable - exec: The command was just executed, w/o returning any variable
parameters	dictionary	Dictionary containing the parameters used for the command execution
ret	(variable)	Variable returned by the command
server	integer	MPI rank of the server process which executed the command request
status	string	Command request status: - timeout: The server assigned to this command requested timeout. - holding queue: This command request is still holding a queue, and has not been sent to the server yet. - request sent: This command has already been sent to a server, and the response has been received.
successful	boolean	Boolean to specify if the command execution was successful or not, where successful means that the command did not throw any exceptions.
traceback	string	When the command is not successful it throws an exception, whose back trace is stored in this parameter.



- `def get_command_response(self,command_request_id_list,block=False,verbose=True)`

This method to retrieve the response for a given command request id/ids specified as a list of integers as returned by `push_command_request` in non-blocking mode.

This method also has a blocking and non-blocking modes, where blocking means that the method will not return until all the responses have been received, and non-blocking will return immediately, giving the responses already available as a list of dictionaries, one per command request, containing the response parameters.

The verbose parameter control the information provided when the command responses have not been received yet. When verbose mode is True it will inform about the servers which are running the command requests whose answer has not been received yet.

- **Command Request Id:** List with ids (integers) of the command response to retrieve
- **Block:** Boolean to control whether command request is executed in blocking or not:
  - ◆ Block=True: It will block until the response from all command is received
  - ◆ Block=False: It will not block, and just return the available responses

**NOTE:** If a command was sent to a server which is not responsive at retrieval time, it will not block the return of this method, but notify of the time-out error instead

- **verbose:** Boolean to control weather information about command request status is posted in non-blocking mode



- **returns:**

- ◆ In non-blocking mode: It will return immediately, giving the responses already available as a list of dictionaries, one per command request, containing the response parameters.

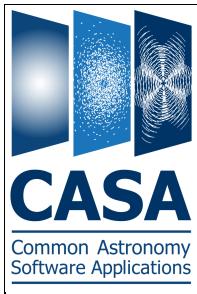
```
CASA <98>: client.push_command_request("a+b",False,[0],{'a':30,'b':40})
Out[98]: [8]
2014-05-19 15:16:38    INFO    casa::MPICommandClient::command_request_queue_service::casa@almahpc02::MPIClient Command request with id# 8 sent to server n# 0
CASA <99>: 2014-05-19 15:16:38  INFO    casa::MPICommandClient::command_response_handler_service::casa@almahpc02::MPIClient    Command request with id 8 successfully handled by server n# 0

CASA <100>: client.push_command_request("time,sleep(a+b)",False,[1],{'a':30,'b':40})
Out[100]: [9]
2014-05-19 15:16:49    INFO    casa::MPICommandClient::command_request_queue_service::casa@almahpc02::MPIClient Command request with id# 9 sent to server n# 1
```

```
CASA <101>: client.get_command_response([8,9],False,True)
2014-05-19 15:17:10    INFO    casa::MPICommandClient::get_command_response::casa@almahpc02::MPIClient Command request with id# 9 assigned to server n# 1, response pending ...
Out[101]:
[{'command': 'a+b',
 'id': 8,
 'mode': 'eval',
 'parameters': {'a': 30, 'b': 40},
 'ret': 70,
 'server': 0,
 'status': 'response received',
 'successful': True,
 'traceback': None}]
```

- ◆ In blocking mode: It will not return until the responses for all the command requests have been received. Then it returns a list of dictionaries, one per command request, containing the response parameters.

```
CASA <102>: client.get_command_response([8,9],True,True)
2014-05-19 15:17:59    INFO    casa::MPICommandClient::command_response_handler_service::casa@almahpc02::MPIClient    Command request with id 9 successfully handled by server n# 1
Out[102]:
[{'command': 'a+b',
 'id': 8,
 'mode': 'eval',
 'parameters': {'a': 30, 'b': 40},
 'ret': 70,
 'server': 0,
 'status': 'response received',
 'successful': True,
 'traceback': None},
 {'command': 'time,sleep(a+b)',
 'id': 9,
 'mode': 'eval',
 'parameters': {'a': 30, 'b': 40},
 'ret': None,
 'server': 1,
 'status': 'response received',
 'successful': True,
 'traceback': None}]
```



## 5.5 MPICommandClient/Server logging

In the current version (CASA 4.3) both sides (client and servers) send the logs to the same log file, and it is possible to inspect the server logs in real time, by using the casalog GUI in the client side. However, the server logs will not appear in the Client terminal directly.

Provided that the logs from all servers are flushed to the same file, the MPI rank of each server has been added to the log origin to facilitate the log analysis. However, this has been done only at the Python level, and in second iteration it will be added to the C++ level. It is a rather trivial step but since it involves compiling CASA with MPI compilers I have decided to postpone this step until we have a proper CASA development environment with MPI enabled.

In future versions the plan is that the logger in the server side sends the logs to the client logger, which in turn can send them to the terminal and/or the log file.

```
CASA <12>: casalogger
-----> casalogger()

Log Messages (almahpc02:/data1/testhpc2/MPI/casapy-20140520-100332.log)
```

File Edit View

Priority Origin Message

Priority	Origin	Message
WARN	flagdata::utils::verify	Some arguments failed to verify!
INFO	...sa@almahpc02:MPIServer-0	Command request with id 1 successfully processed in eval mode, sending to server n# 0
INFO	...casa@almahpc02:MPIClient	Command request with id 1 successfully handled by server n# 0
INFO	...casa@almahpc02:MPIClient	Command request with id# 2 sent to server n# 0
INFO	...sa@almahpc02:MPIServer-0	Received command request msg: myimager = imager()
INFO	...sa@almahpc02:MPIServer-0	Going to execute command request with id# 2 as a statement via exec
INFO	...sa@almahpc02:MPIServer-0	Command request with id 2 successfully processed in exec mode, sending to server n# 0
INFO	...casa@almahpc02:MPIClient	Command request with id 2 successfully handled by server n# 0
INFO	...casa@almahpc02:MPIClient	Command request with id# 3 sent to server n# 0
INFO	...sa@almahpc02:MPIServer-0	Received command request msg: myimager.open(vis)
INFO	...sa@almahpc02:MPIServer-0	Going to evaluate command request with id# 3 as an expression via eval
SEVERE	...sa@almahpc02:MPIServer-0	Exception executing command request via eval: Traceback (most recent call last):
SEVERE	...a@almahpc02:MPIServer-0+	File "/data1/jagonzal/mpi4py/linux_64b/python/2.7/mpi4casa/MPICommander.py", line 1, in <module>
SEVERE	...a@almahpc02:MPIServer-0+	command_response['ret'] = eval(command_request['command'])
SEVERE	...a@almahpc02:MPIServer-0+	File "<string>", line 1, in <module>
SEVERE	...a@almahpc02:MPIServer-0+	File "/data1/jagonzal/mpi4py/linux_64b/python/2.7/__casac__/image.py", line 1, in <module>
SEVERE	...a@almahpc02:MPIServer-0+	return _imager.imager_open(self, *args, **kwargs)
SEVERE	...a@almahpc02:MPIServer-0+	RuntimeError: Table /data1/testhpc2/MPI/fileNotFound.ms does not exist
INFO	...sa@almahpc02:MPIServer-0	Command request with id 3 successfully processed in eval mode, sending to server n# 0
INFO	...casa@almahpc02:MPIClient	Command request with id 3 failed in server n# 0 with traceback Traceback (most recent call last):
INFO	...asa@almahpc02:MPIClient+	File "/data1/jagonzal/mpi4py/linux_64b/python/2.7/mpi4casa/MPICommander.py", line 1, in <module>
INFO	...asa@almahpc02:MPIClient+	command_response['ret'] = eval(command_request['command'])

Insert Message:     Lock scroll



## 5.6 MPICommandClient/Server error handling

MPIClient/Server are ready to handle any python exception that occurs in the server, and report it back to the client side, properly formatted.

When a python/CASA command is not successful (i.e. and exception is thrown), then the command response sets the Boolean field 'successful' to False, and the string field 'traceback' contains the python trace-back already formatted (using the python module traceback).

**NOTE:** Please notice that in terms of Python a command execution is considered not successful only when it throws an exception. In terms of CASA, some tasks return a Boolean False when the call is not successful, but w/o throwing any exception. In this case, the return variable 'ret' of the command response will contain a Boolean False, but the Boolean field 'successful' will be set to True.

- ex 1.: exception thrown when opening a tool

```
parameters = {'vis':'fileNotFound.ms'}
response = client.push_command_request('myimagertool.open(vis)',True,[0],parameters)

CASA <11>: if not response[0]['successful']:
....:     print response[0]['traceback']
....:
....:
Traceback (most recent call last):
  File "/data1/jagonzal/mpi4py/linux_64b/python/2.7/mpi4casa/MPICommandserver.py", line
145, in __command_request_handler_service
    command_response['ret'] = eval(command_request['command'])
  File "<string>", line 1, in <module>
  File "/data1/jagonzal/mpi4py/linux_64b/python/2.7/_casac__imager.py", line 1932, in
open
    return _imager.imager_open(self, *args, **kwargs)
RuntimeError: Table /data1/testhpc2/MPI/fileNotFound.ms does not exist
```

```
CASA <10>: response = client.push_command_request('myimagertool.open(vis)',True,[0],parameters={'vis':'fileNotFound.ms'})
2014-05-20 10:07:37 INFO casa::MPICommandClient::command_request_queue_service::casa@almahpc02:MPIClient Command request with id# 4 sent to server n# 0
2014-05-20 10:07:37 INFO casa::MPICommandClient::command_response_handler_service::casa@almahpc02:MPIClient Command request with id 4 failed in server n# 0 with traceback Traceback (most recent call last):
2014-05-20 10:07:37 INFO casa::MPICommandClient::command_response_handler_service::casa@almahpc02:MPIClient   File "/data1/jagonzal/mpi4py/linux_64b/python/2.7/mpi4casa/MPICommandServer.py", line 145, in __
....:     command_response['ret'] = eval(command_request['command'])
....:   File "<string>", line 1, in <module>
....:   File "/data1/jagonzal/mpi4py/linux_64b/python/2.7/_casac__imager.py", line 1932, in open
....:     return _imager.imager_open(self, *args, **kwargs)
....: RuntimeError: Table /data1/testhpc2/MPI/fileNotFound.ms does not exist

CASA <11>: if not response[0]['successful']:
....:     print response[0]['traceback']
....:
....:
....:
Traceback (most recent call last):
  File "/data1/jagonzal/mpi4py/linux_64b/python/2.7/mpi4casa/MPICommandServer.py", line 145, in __command_request_handler_service
    command_response['ret'] = eval(command_request['command'])
  File "<string>", line 1, in <module>
  File "/data1/jagonzal/mpi4py/linux_64b/python/2.7/_casac__imager.py", line 1932, in open
    return _imager.imager_open(self, *args, **kwargs)
RuntimeError: Table /data1/testhpc2/MPI/fileNotFound.ms does not exist
```



## 5.7 MPICommandClient basic python examples

- **ex 1:** Blocking mode, string parameters, undefined target server

```
command_response_list = client.push_command_request("I+I",True,None)
```

```
CASA <5>: command_response_list = client.push_command_request("1+1",True,None)
2014-05-19 14:02:32      INFO  casa::MPICommandClient::command_request_queue_service::casa@almahpc02:MPIClient Command request with id# 1 sent to server n# 2
2014-05-19 14:02:33      INFO  casa::MPICommandClient::command_response_handler_service::casa@almahpc02:MPIClient      Command request with id 1 successfully handled by server n# 2

CASA <6>: command_response_list
Out[6]:
[{'command': '1+1',
 'id': 1,
 'mode': 'eval',
 'parameters': None,
 'ret': 2,
 'server': 2,
 'status': 'response received',
 'successful': True,
 'traceback': None}]
```

- **ex 2:** Blocking mode, string parameters, defined target server

```
command_response_list = client.push_command_request("I+I",True,[0])
```

```
CASA <7>: command_response_list = client.push_command_request("1+1",True,[0])
2014-05-19 14:02:51      INFO  casa::MPICommandClient::command_request_queue_service::casa@almahpc02:MPIClient Command request with id# 2 sent to server n# 0
2014-05-19 14:02:51      INFO  casa::MPICommandClient::command_response_handler_service::casa@almahpc02:MPIClient      Command request with id 2 successfully handled by server n# 0

CASA <8>: command_response_list
Out[8]:
[{'command': '1+1',
 'id': 2,
 'mode': 'eval',
 'parameters': None,
 'ret': 2,
 'server': 0,
 'status': 'response received',
 'successful': True,
 'traceback': None}]
```



- **ex 3:** Blocking mode, string parameters, multiple target server

```
command_response_list = client.push_command_request("I+I",True,[0,1])
```

```
CASA <9>: command_response_list = client.push_command_request("1+1",True,[0,1])
2014-05-19 14:03:12    INFO    casa::MPICommandClient:::command_request_queue_service::casa@almahpc02:MPIClient Command request with id# 3 sent to server n# 0
2014-05-19 14:03:12    INFO    casa::MPICommandClient:::command_request_queue_service::casa@almahpc02:MPIClient Command request with id# 4 sent to server n# 1
2014-05-19 14:03:13    INFO    casa::MPICommandClient:::command_response_handler_service::casa@almahpc02:MPIClient      Command request with id 4 successfully handled by server n# 1
2014-05-19 14:03:13    INFO    casa::MPICommandClient:::command_response_handler_service::casa@almahpc02:MPIClient      Command request with id 3 successfully handled by server n# 0

CASA <10>: command_response_list
Out[10]:
[{'command': '1+1',
 'id': 3,
 'mode': 'eval',
 'parameters': None,
 'ret': 2,
 'server': 0,
 'status': 'response received',
 'successful': True,
 'traceback': None},
 {'command': '1+1',
 'id': 4,
 'mode': 'eval',
 'parameters': None,
 'ret': 2,
 'server': 1,
 'status': 'response received',
 'successful': True,
 'traceback': None}]
```

- **ex 4:** Non-Blocking mode, string parameters, undefined target server

```
command_request_id_list = self.client.push_command_request("I+I",False,None)
# Try to get responses before time in non-blocking mode
command_response_list = client.get_command_response(command_request_id_list,False,True)
# Get response in blocking mode
command_response_list = client.get_command_response(command_request_id_list,True,True)
```

```
CASA <11>: command_request_id_list = client.push_command_request("1+1",False,None)
2014-05-19 14:03:27    INFO    casa::MPICommandClient:::command_request_queue_service::casa@almahpc02:MPIClient Command request with id# 5 sent to server n# 2
CASA <12>: 2014-05-19 14:03:28 INFO    casa::MPICommandClient:::command_response_handler_service::casa@almahpc02:MPIClient      Command request with id 5 successfully handled by server n# 2

CASA <13>: command_response_list = client.get_command_response(command_request_id_list,False,True)
CASA <14>: command_response_list = client.get_command_response(command_request_id_list,True,True)

CASA <15>: command_response_list
Out[15]:
[{'command': '1+1',
 'id': 5,
 'mode': 'eval',
 'parameters': None,
 'ret': 2,
 'server': 2,
 'status': 'response received',
 'successful': True,
 'traceback': None}]
```



- **ex 5:** Non-Blocking mode, string parameters, defined target server

```
command_request_id_list = self.client.push_command_request("1+1",False,[0])
# Try to get responses before time in non-blocking mode
command_response_list = self.client.get_command_response(command_request_id_list,False,True)
# Get response in blocking mode
command_response_list
```

```
CASA <16>: command_request_id_list = client.push_command_request("1+1",False,[0])
2014-05-19 14:05:36    INFO    casa::MPICommandClient::command_request_queue_service::casa@almahpc02:MPIClient Command request with id# 6 sent to server n# 0

CASA <17>: 2014-05-19 14:05:36  INFO    casa::MPICommandClient::command_response_handler_service::casa@almahpc02:MPIClient      Command request with id 6 successfully handled by server n# 0

CASA <18>: command_response_list = client.get_command_response(command_request_id_list,False,True)
CASA <19>: command_response_list = client.get_command_response(command_request_id_list,True,True)

CASA <20>: command_response_list
Out[20]:
[{'command': '1+1',
 'id': 6,
 'mode': 'eval',
 'parameters': None,
 'ret': 2,
 'server': 0,
 'status': 'response received',
 'successful': True,
 'traceback': None}]
```

▪ **ex 6:** Non-Blocking mode, string parameters, multiple target servers

```
command_request_id_list = self.client.push_command_request("1+1",False,[0,1])
# Try to get responses before time in non-blocking mode
command_response_list = self.client.get_command_response(command_request_id_list,False,True)
# Get response in blocking mode
command_response_list = self.client.get_command_response(command_request_id_list,True,True)
```

```
CASA <21>: command_request_id_list = client.push_command_request("1+1",False,[0,1])
```

```
2014-05-19 14:06:13    INFO    casa::MPICommandClient::command_request_queue_service::casa@almahpc02:MPIClient Command request with id# 7 sent to server n# 0
CASA <22>: 2014-05-19 14:06:13  INFO    casa::MPICommandClient::command_request_queue_service::casa@almahpc02:MPIClient Command request with id# 8 sent to server n# 1
2014-05-19 14:06:14    INFO    casa::MPICommandClient::command_response_handler_service::casa@almahpc02:MPIClient      Command request with id 8 successfully handled by server n# 1
2014-05-19 14:06:14    INFO    casa::MPICommandClient::command_response_handler_service::casa@almahpc02:MPIClient      Command request with id 7 successfully handled by server n# 0
CASA <22>: command_response_list = client.get_command_response(command_request_id_list,False,True)
```

```
CASA <23>: command_response_list = client.get_command_response(command_request_id_list,True,True)
```

```
CASA <24>: command_response_list
```

```
Out[24]:
[{'command': '1+1',
 'id': 7,
 'mode': 'eval',
 'parameters': None,
 'ret': 2,
 'server': 0,
 'status': 'response received',
 'successful': True,
 'traceback': None},
 {'command': '1+1',
 'id': 8,
 'mode': 'eval',
 'parameters': None,
 'ret': 2,
 'server': 1,
 'status': 'response received',
 'successful': True,
 'traceback': None}]
```

▪ **ex 7:** Blocking mode, dictionary parameters, undefined target server

```
command_response_list = self.client.push_command_request("a+b",True,None,{"a":1,'b':2})
```

```
CASA <25>: command_response_list = client.push_command_request("a+b",True,None,{"a":1,'b':2})
```

```
2014-05-19 14:06:56    INFO    casa::MPICommandClient::command_request_queue_service::casa@almahpc02:MPIClient Command request with id# 9 sent to server n# 2
2014-05-19 14:06:56    INFO    casa::MPICommandClient::command_response_handler_service::casa@almahpc02:MPIClient      Command request with id 9 successfully handled by server n# 2
```

```
CASA <26>: command_response_list
```

```
Out[26]:
[{'command': 'a+b',
 'id': 9,
 'mode': 'eval',
 'parameters': {'a': 1, 'b': 2},
 'ret': 3,
 'server': 2,
 'status': 'response received',
 'successful': True,
 'traceback': None}]
```

▪ **ex 8:** Blocking mode, dictionary parameters, defined target server

```
command_response_list = self.client.push_command_request("a+b",True,[0],{'a':1,'b':2})
```

```
CASA <27>: command_response_list = client.push_command_request("a+b",True,[0],{'a':1,'b':2})
2014-05-19 14:07:13    INFO    casa::MPICommandClient::command_request_queue_service::casa@almahpc02::MPIClient Command request with id# 10 sent to server n# 0
2014-05-19 14:07:14    INFO    casa::MPICommandClient::command_response_handler_service::casa@almahpc02::MPIClient    Command request with id 10 successfully handled by server n# 0

CASA <28>: command_response_list
Out[28]:
[{'command': 'a+b',
 'id': 10,
 'mode': 'eval',
 'parameters': {'a': 1, 'b': 2},
 'ret': 3,
 'server': 0,
 'status': 'response received',
 'successful': True,
 'traceback': None}]
```

▪ **ex 9:** Blocking mode, dictionary parameters, multiple target server

```
command_response_list = self.client.push_command_request("a+b",True,[0,1],{'a':1,'b':2})
```

```
CASA <29>: command_response_list = client.push_command_request("a+b",True,[0,1],{'a':1,'b':2})
2014-05-19 14:07:27    INFO    casa::MPICommandClient::command_request_queue_service::casa@almahpc02::MPIClient Command request with id# 11 sent to server n# 0
2014-05-19 14:07:27    INFO    casa::MPICommandClient::command_request_queue_service::casa@almahpc02::MPIClient Command request with id# 12 sent to server n# 1
2014-05-19 14:07:28    INFO    casa::MPICommandClient::command_response_handler_service::casa@almahpc02::MPIClient    Command request with id 11 successfully handled by server n# 0
2014-05-19 14:07:28    INFO    casa::MPICommandClient::command_response_handler_service::casa@almahpc02::MPIClient    Command request with id 12 successfully handled by server n# 1

CASA <30>: command_response_list
Out[30]:
[{'command': 'a+b',
 'id': 11,
 'mode': 'eval',
 'parameters': {'a': 1, 'b': 2},
 'ret': 3,
 'server': 0,
 'status': 'response received',
 'successful': True,
 'traceback': None},
 {'command': 'a+b',
 'id': 12,
 'mode': 'eval',
 'parameters': {'a': 1, 'b': 2},
 'ret': 3,
 'server': 1,
 'status': 'response received',
 'successful': True,
 'traceback': None}]
```



- **ex 10:** Non-Blocking mode, dictionary parameters, undefined target server

```
command_request_id_list = self.client.push_command_request("a+b",False,None,{‘a’:1,’b’:2})  
# Try to get responses before time in non-blocking mode  
command_response_list = self.client.get_command_response(command_request_id_list,False,True)  
# Get response in blocking mode  
command_response_list = self.client.get_command_response(command_request_id_list,True,True)
```

```
CASA <31>; command_request_id_list = client.push_command_request("a+b",False,None,{‘a’:1,’b’:2})  
CASA <32>; 2014-05-19 14:07:43 INFO casa::MPICommandClient::command_request_queue_service::casa@almahpc02:MPIClient Command request with id# 13 sent to server n# 2  
2014-05-19 14:07:44 INFO casa::MPICommandClient::command_response_handler_service::casa@almahpc02:MPIClient Command request with id 13 successfully handled by server n# 2  
  
CASA <33>; command_response_list = client.get_command_response(command_request_id_list,False,True)  
CASA <34>; command_response_list = client.get_command_response(command_request_id_list,True,True)  
  
CASA <35>; command_response_list  
Out[35]:  
[{'command': 'a+b',  
 'id': 13,  
 'mode': 'eval',  
 'parameters': {'a': 1, 'b': 2},  
 'ret': 3,  
 'server': 2,  
 'status': 'response received',  
 'successful': True,  
 'traceback': None}]
```

- **ex 11:** Non-Blocking mode, dictionary parameters, defined target server

```
command_request_id_list = self.client.push_command_request("a+b",False,[0],{‘a’:1,’b’:2})  
# Try to get responses before time in non-blocking mode  
command_response_list = self.client.get_command_response(command_request_id_list,False,True)  
# Get response in blocking mode  
command_response_list = self.client.get_command_response(command_request_id_list,True,True)
```

```
CASA <36>; command_request_id_list = client.push_command_request("a+b",False,[0],{‘a’:1,’b’:2})  
2014-05-19 14:08:19 INFO casa::MPICommandClient::command_request_queue_service::casa@almahpc02:MPIClient Command request with id# 14 sent to server n# 0  
CASA <37>; 2014-05-19 14:08:20 INFO casa::MPICommandClient::command_response_handler_service::casa@almahpc02:MPIClient Command request with id 14 successfully handled by server n# 0  
CASA <37>; command_response_list = client.get_command_response(command_request_id_list,False,True)  
  
CASA <38>; command_response_list = client.get_command_response(command_request_id_list,True,True)  
  
CASA <39>; command_response_list  
Out[39]:  
[{'command': 'a+b',  
 'id': 14,  
 'mode': 'eval',  
 'parameters': {'a': 1, 'b': 2},  
 'ret': 3,  
 'server': 0,  
 'status': 'response received',  
 'successful': True,  
 'traceback': None}]
```

▪ **ex 12:** Non-Blocking mode, dictionary parameters, multiple target server

```
command_request_id_list = self.client.push_command_request("a+b",False,[0,1],{'a':1,'b':2})  
# Try to get responses before time in non-blocking mode  
command_response_list = self.client.get_command_response(command_request_id_list,False,True)  
# Get response in blocking mode  
command_response_list = self.client.get_command_response(command_request_id_list,True,True)
```

```
CASA <40>: command_request_id_list = client.push_command_request("a+b",False,[0,1],{'a':1,'b':2})
```

```
2014-05-19 14:08:48    INFO    casa::MPICommandClient::command_request_queue_service::casa@almahpc02:MPIClient Command request with id# 15 sent to server n# 0  
CASA <41>: 2014-05-19 14:08:48    INFO    casa::MPICommandClient::command_request_queue_service::casa@almahpc02:MPIClient Command request with id# 16 sent to server n# 1  
2014-05-19 14:08:49    INFO    casa::MPICommandClient::command_response_handler_service::casa@almahpc02:MPIClient    Command request with id 16 successfully handled by server n# 1  
2014-05-19 14:08:49    INFO    casa::MPICommandClient::command_response_handler_service::casa@almahpc02:MPIClient    Command request with id 15 successfully handled by server n# 0  
CASA <41>: command_response_list = client.get_command_response(command_request_id_list,False,True)
```

```
CASA <42>: command_response_list = client.get_command_response(command_request_id_list,True,True)
```

```
CASA <43>: command_response_list  
Out[43]:  
[{'command': 'a+b',  
'id': 15,  
'mode': 'eval',  
'parameters': {'a': 1, 'b': 2},  
'ret': 3,  
'server': 0,  
'status': 'response received',  
'successful': True,  
'traceback': None},  
{'command': 'a+b',  
'id': 16,  
'mode': 'eval',  
'parameters': {'a': 1, 'b': 2},  
'ret': 3,  
'server': 1,  
'status': 'response received',  
'successful': True,  
'traceback': None}]
```



## 5.1 MPICommandClient CASA tasks examples

Executing a CASA task command does not differ at all from executing a basic python command, but for the sake of completeness I have prepared some basic examples here:

**NOTE:** mpi4py is already integrated with ParallelTaskHelper, so it is not necessary to do anything (in particular initialize MPICommandClient) in order to run tasks in parallel. ParallelTaskHelper re-uses or initializes the MPICommandClient singleton if necessary.

- ex 1.: flagdata rflag, string parameters, blocking mode, defined server

```
client.push_command_request("flagdata(vis='Four_ants_3C286.ms',mode='summary')",True,[1])
```

```
CASA <5>; parameters={}
CASA <6>; parameters['vis']="Four_ants_3C286.ms"
CASA <7>; parameters['mode']='rflag'
CASA <8>; parameters['backup']=False
CASA <9>; client.push_command_request('flagdata()',True,[0],parameters)
2014-05-19 17:11:12    INFO    casa::MPICommandClient::command_request_queue_service::casa@almahpc02:MPIClient Command request with id# 1 sent to server n# 0
2014-05-19 17:11:18    INFO    casa::MPICommandClient::command_response_handler_service::casa@almahpc02:MPIClient    Command request with id 1 successfully handled by server n# 0
Out[9]:
[{'command': 'flagdata()',
 'id': 1,
 'mode': 'eval',
 'parameters': {'backup': False,
                'mode': 'rflag',
                'vis': 'Four_ants_3C286.ms'},
 'ret': {},
 'server': 0,
 'status': 'response received',
 'successful': True,
 'traceback': None}]
```



- ex 2.: flagdata summary, dict parameters, blocking mode, defined server

```
parameters = {'vis':'Four_ants_3C286.ms','mode':'rflag','flagbackup':False}
client.push_command_request('flagdata()',True,[0],parameters)
```

```
CASA <5>: parameters={}
CASA <6>: parameters['vis']="Four_ants_3C286.ms"
CASA <7>: parameters['mode']='rflag'
CASA <8>: parameters['backup']=False
CASA <9>: client.push_command_request('flagdata()',True,[0],parameters)
2014-05-19 17:11:12    INFO    casa::MPICommandClient::command_request_queue_service::casa@almahpc02:MPIClient Command request with id# 1 sent to server n# 0
2014-05-19 17:11:18    INFO    casa::MPICommandClient::command_response_handler_service::casa@almahpc02:MPIClient    Command request with id 1 successfully handled by server n# 0
Out[9]:
[{'command': 'flagdata',
 'id': 1,
 'mode': 'eval',
 'parameters': {'backup': False,
                'mode': 'rflag',
                'vis': 'Four_ants_3C286.ms'},
 'ret': {},
 'server': 0,
 'status': 'response received',
 'successful': True,
 'traceback': None}]
```



- ex 3.: create an imager tool in each remote server, blocking mode

```
client.push_command_request('myimager=imager()',True,[0,1,2])
```

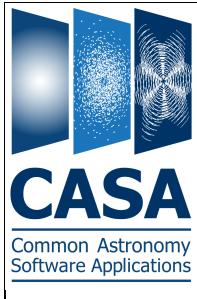
```
CASA <5> client.push_command_request("myimager=imager()",True,[0,1,2])
2014-05-19 17:21:02    INFO  casa::MPICommandClient::command_request_queue_service::casa@almahpc02;MPIClient Command request with id# 1 sent to server n# 0
2014-05-19 17:21:02    INFO  casa::MPICommandClient::command_request_queue_service::casa@almahpc02;MPIClient Command request with id# 2 sent to server n# 1
2014-05-19 17:21:02    INFO  casa::MPICommandClient::command_request_queue_service::casa@almahpc02;MPIClient Command request with id# 3 sent to server n# 2
2014-05-19 17:21:02    INFO  casa::MPICommandClient::command_response_handler_service::casa@almahpc02;MPIClient      Command request with id 2 successfully handled by server n# 1
2014-05-19 17:21:02    INFO  casa::MPICommandClient::command_response_handler_service::casa@almahpc02;MPIClient      Command request with id 1 successfully handled by server n# 0
2014-05-19 17:21:02    INFO  casa::MPICommandClient::command_response_handler_service::casa@almahpc02;MPIClient      Command request with id 3 successfully handled by server n# 2
Out[5]:
[{'command': 'myimager=imager()',
 'id': 1,
 'mode': 'exec',
 'parameters': None,
 'ret': None,
 'server': 0,
 'status': 'response received',
 'successful': True,
 'traceback': None},
 {'command': 'myimager=imager()',
 'id': 2,
 'mode': 'exec',
 'parameters': None,
 'ret': None,
 'server': 1,
 'status': 'response received',
 'successful': True,
 'traceback': None},
 {'command': 'myimager=imager()',
 'id': 3,
 'mode': 'exec',
 'parameters': None,
 'ret': None,
 'server': 2,
 'status': 'response received',
 'successful': True,
 'traceback': None}]
```



- ex 4.: open newly created imager tool in each remote server, dict parameters, blocking mode

```
parameters={'vis':'Four_ants_3C286.ms'}
client.push_command_request('myimager.open(vis)',True,[0,1,2],parameters)
```

```
CASA <8>; client.push_command_request("myimager.open(vis)",True,[0,1,2],parameters={'vis':'Four_ants_3C286.ms'})
2014-05-19 17:22:25    INFO  casa::MPICommandClient::command_request_queue_service::casa@almahpc02:MPIClient Command request with id# 4 sent to server n# 0
2014-05-19 17:22:25    INFO  casa::MPICommandClient::command_request_queue_service::casa@almahpc02:MPIClient Command request with id# 5 sent to server n# 1
2014-05-19 17:22:25    INFO  casa::MPICommandClient::command_request_queue_service::casa@almahpc02:MPIClient Command request with id# 6 sent to server n# 2
2014-05-19 17:22:25    INFO  casa::MPICommandClient::command_response_handler_service::casa@almahpc02:MPIClient      Command request with id 5 successfully handled by server n# 1
2014-05-19 17:22:25    INFO  casa::MPICommandClient::command_response_handler_service::casa@almahpc02:MPIClient      Command request with id 4 successfully handled by server n# 0
2014-05-19 17:22:25    INFO  casa::MPICommandClient::command_response_handler_service::casa@almahpc02:MPIClient      Command request with id 6 successfully handled by server n# 2
Out[8]:
[{'command': 'myimager.open(vis)',
 'id': 4,
 'mode': 'eval',
 'parameters': {'vis': 'Four_ants_3C286.ms'},
 'ret': True,
 'server': 0,
 'status': 'response received',
 'successful': True,
 'traceback': None},
 {'command': 'myimager.open(vis)',
 'id': 5,
 'mode': 'eval',
 'parameters': {'vis': 'Four_ants_3C286.ms'},
 'ret': True,
 'server': 1,
 'status': 'response received',
 'successful': True,
 'traceback': None},
 {'command': 'myimager.open(vis)',
 'id': 6,
 'mode': 'eval',
 'parameters': {'vis': 'Four_ants_3C286.ms'},
 'ret': True,
 'server': 2,
 'status': 'response received',
 'successful': True,
 'traceback': None}]
```



## 5.2 mpi4casa unit test suit

mpi4casa has a complete unit test suit, with more than 60 tests, which can be launched using CASA's unit test tools, specifying a configurable number of parallel processors, or even specifying a distributed environment via hostfile.

```
mpirun -n 4 casapy --noLogger --log2term -c runUnitTest.py test_mpi4casa
```

```
mpirun -n 4 -hostfile host casapy --noLogger --log2term -c runUnitTest.py test_mpi4casa
```

The unit test suit covers all combinations of MPICommandClient/Server usage:

- blocking and non-blocking mode
- execution mode
  - with return variable (eval)
  - w/o return variable (exec)
- target server specification
  - undefined server
  - defined server
  - multiple servers
  - busy/idle server
- parameters specification
  - Via command string
  - Via parameter dictionary
- error handling cases:
  - ImportError
  - NameError
  - ZeroDivisionError
  - TypeError
  - IndexError
  - KeyError
- Server timeout handling
  - Server timeout
  - Assignment to timeout server
  - Timeout recovery
- Singleton behavior
  - Invalid MPICommandClient/Server instantiation
  - MPICommandClient destruction

It also covers integration via ParallelTaskHelper with the following tasks:

- flagdata
- applycal
- uvcontsub
- setjy